

Lecture Slides - 4

Classification & Prediction

Classification and Prediction

- What is classification? What is prediction?
- Issues regarding classification and prediction
- Classification techniques
- Prediction
- Accuracy and error measures
- Ensemble methods
- Model selection
- Summary

Classification vs. Prediction

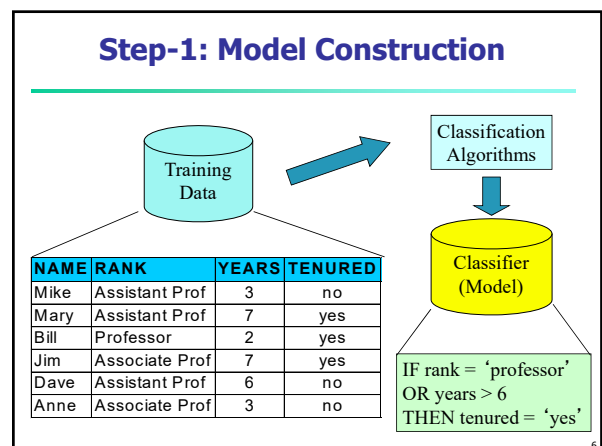
- Classification**
 - Predicts categorical class labels (discrete or nominal)
 - Classifies data (constructs a model) based on the training set and the values (**class labels**) in a classifying attribute and uses it in classifying new data
- Prediction**
 - Models continuous-valued functions, i.e., predicts unknown or missing values
- Typical applications**
 - Credit approval
 - Target marketing
 - Medical diagnosis
 - Fraud detection

Classification: A Mathematical Mapping

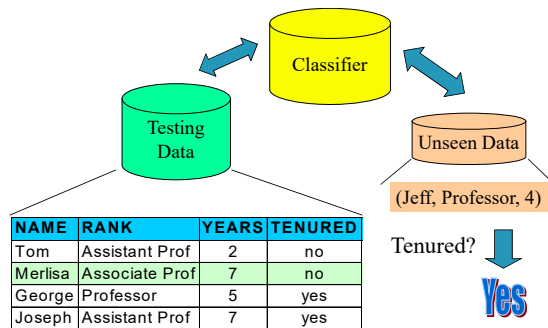
- Classification:**
 - predicts categorical class labels
- Mathematically**
 - $x \in X = \mathbb{R}^n, y \in Y = \{+1, -1\}$
 - We want a function $f: X \rightarrow Y$
- E.g., Personal homepage classification**
 - $x_i = (x_1, x_2, x_3, \dots), y_i = +1 \text{ or } -1$
 - x_1 : # of a word "homepage"
 - x_2 : # of a word "welcome"

Classification — A Two-Step Process

- Model construction:** describing a set of predetermined classes
 - Each tuple/sample is assumed to belong to a predefined class, as determined by the **class label attribute**
 - The set of tuples used for model construction is **training set**
 - The model is represented as classification rules, decision trees, or mathematical formulae
- Model usage:** for classifying future or unknown objects
 - Estimate accuracy** of the model
 - The known label of test sample is compared with the classified result from the model
 - Accuracy rate is the percentage of test set samples that are correctly classified by the model
 - If the accuracy is acceptable, use the model to **classify data** tuples whose class labels are not known



Step-2: Using the Model in Prediction



7

Supervised vs. Unsupervised Learning

- **Supervised learning (classification)**
 - Supervision: The training data (observations, measurements, etc.) are accompanied by labels indicating the class of the observations
 - New data is classified based on the training set
- **Unsupervised learning (clustering)**
 - The class labels of training data is unknown
 - Given a set of measurements, observations, etc. the aim is to establish the existence of classes or clusters in the data

8

Issues: Data Preparation

- **Data cleaning**
 - Preprocess data in order to reduce noise and handle missing values
- **Relevance analysis (feature selection)**
 - Remove the irrelevant or redundant attributes
- **Data transformation**
 - Generalize and/or normalize data

9

Issues: Evaluating Classification Methods

- **Accuracy**
 - Classifier accuracy: predicting class label
 - Predictor accuracy: guessing value of predicted attributes
- **Speed**
 - Time to construct the model (training time)
 - Time to use the model (classification/prediction time)
- **Robustness:**
 - Handling noise and missing values
- **Scalability:**
 - Efficiency in disk-resident databases
- **Interpretability**
 - Understanding and insight provided by the model
- **Other measures**, e.g., goodness of rules, such as decision tree size or compactness of classification rules


10

Classification Techniques

- **Well-known methods:**
 - Decision Trees
 - Bayesian Classification
 - Bayesian Belief Networks
- **Besides, we can also use**
 - Instance-Based Methods
 - k -nearest neighbor (k -NN) approach
 - Neural Networks
 - Support Vector Machine

Classification by DT Induction

- **Decision tree**
 - A flow-chart-like tree structure
 - Internal node denotes a test on an attribute
 - Branch represents an outcome of the test
 - Leaf nodes represent class labels or class distribution
- **Decision tree generation consists of two phases:**
 - Tree construction
 - Tree pruning
 - Identify and remove branches that reflect noise or outliers
- Use of decision tree: **Classifying unknown samples**



Decision Tree Induction: Training Dataset

age	income	student	credit_rating	buys_computer
<=30	high	no	fair	no
<=30	high	no	excellent	no
31...40	high	no	fair	yes
>40	medium	no	fair	yes
>40	low	yes	fair	yes
>40	low	yes	excellent	no
31...40	low	yes	excellent	yes
<=30	medium	no	fair	no
<=30	low	yes	fair	yes
>40	medium	yes	fair	yes
<=30	medium	yes	excellent	yes
31...40	medium	no	excellent	yes
31...40	high	yes	fair	yes
>40	medium	no	excellent	no

Algorithm for Decision Tree Induction

- Basic algorithm (a greedy algorithm)
 - Tree is constructed in a **top-down recursive divide-and-conquer manner**
 - At start, all the training examples are at the root
 - Attributes are categorical (if continuous-valued, they are discretized in advance)
 - Examples are partitioned recursively based on selected attributes
 - Test attributes are selected on the basis of a heuristic or statistical measure (e.g., **information gain**)
- Conditions for stopping partitioning
 - All samples for a given node belong to the same class
 - There are no remaining attributes for further partitioning – **majority voting** is employed for classifying the leaf
 - There are no samples left

Decision Tree Algorithms (ID3, C4.5, C5.0/See5)

- ID3 (Iterative Dichotomiser 3)**: A decision tree algorithm invented by Ross Quinlan (RQ)
- C4.5**: An improvement of ID3 by RQ
 - Handles both continuous & discrete attributes
 - Handles training data with missing attributes
 - Prunes tree after creation
- C5.0/See5**: Another improvement by RQ
- C5.0 - Unix/Linux; See5 - Windows
 - Speed
 - Memory usage
 - Smaller decision trees

Attribute Selection Measure: Information Gain

- Select the attribute with the **highest information gain**
- Let p_i be the probability that an arbitrary tuple in D belongs to class C_i , estimated by $|C_{i,D}|/|D|$
- Expected information (entropy)** needed to classify a tuple in D:

$$Info(D) = - \sum_{i=1}^m p_i \log_2(p_i)$$

- Information needed** (after using A to split D into v partitions) to classify D:

$$Info_A(D) = \sum_{j=1}^v \frac{|D_j|}{|D|} \times I(D_j)$$

- Information gained** by branching on attribute A

$$Gain(A) = Info(D) - Info_A(D)$$

Attribute Selection Measure: Information Gain

- Entropy H(D)**:
 - Measure of the amount of uncertainty in the dataset D
- $$H(D) = - \sum_{c \in C} p(c) \log_2 p(c)$$
- Where, D: The dataset for which entropy is calculated
 - C: The set of classes in D
 - $p(c)$: The proportion of the number of elements in class c to the number of elements in D
- When $H(D) = 0$, the dataset D is perfectly classified**
 - i.e. all elements in D are of the same class

Attribute Selection Measure: Information Gain

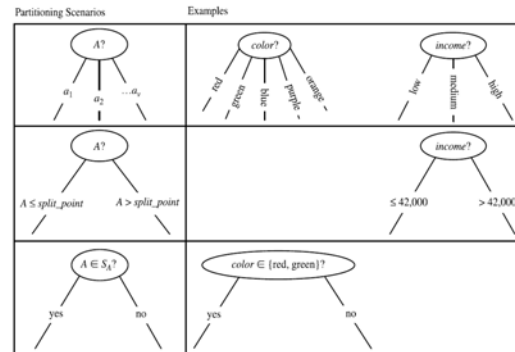
- Information Gain IG(A)**:
 - Measure of the difference in entropy from before to after the dataset D is split on an attribute A.
 - i.e., how much uncertainty in D was reduced after splitting D on attribute A.
- $$IG(A, D) = H(D) - \sum_{s \in S} p(s) H(s)$$
- Where, H(D): Entropy of the dataset D
 - S: The subsets created from splitting D by attribute A
 - $p(s)$: The proportion of the number of elements in s to the number of elements in D
 - H(s): Entropy of the subset s
 - The attribute with the largest information gain is used to split the dataset D.**

Computing Information-Gain for Continuous-Valued Attributes

- Let attribute A be a continuous-valued attribute
- Must determine the **best split point** for A
 - Sort the value A in increasing order
 - Typically, the midpoint between each pair of adjacent values is considered as a possible **split point**
 - $(a_i + a_{i+1})/2$ is the midpoint between the values of a_i and a_{i+1}
 - The point with the **minimum expected information requirement** for A is selected as the split-point for A
- Split:
 - D1 is the set of tuples in D satisfying $A \leq \text{split-point}$, and
 - D2 is the set of tuples in D satisfying $A > \text{split-point}$

19

Splitting Approaches in DT



20

Attribute Selection: Information Gain

- Class P: buys_computer = "yes"
- Class N: buys_computer = "no"

$$\text{Info}(D) = I(9,5) = -\frac{9}{14} \log_2 \left(\frac{9}{14} \right) - \frac{5}{14} \log_2 \left(\frac{5}{14} \right) = 0.940$$

age	p _i	n _i	I(p _i , n _i)
<=30	2	3	0.971
31...40	4	0	0
>40	3	2	0.971

$\frac{5}{14} I(2,3)$ means "age <=30" has 5 out of 14 samples, with 2 yes' es and 3 no' s. Hence

$$\text{Gain}(\text{age}) = \text{Info}(D) - \text{Info}_{\text{age}}(D) = 0.246$$

Similarly,


$$\text{Gain}(\text{income}) = 0.029$$

$$\text{Gain}(\text{student}) = 0.151$$

$$\text{Gain}(\text{credit_rating}) = 0.048$$

age	income	student	credit_rating	buys_computer
<=30	high	no	fair	no
<=30	high	no	excellent	no
31..40	high	no	fair	yes
>40	medium	no	fair	yes
>40	low	yes	fair	yes
>40	low	yes	excellent	no
31..40	low	yes	excellent	yes
<=30	medium	no	fair	no
<=30	low	yes	fair	yes
>40	medium	yes	fair	yes
<=30	medium	yes	excellent	yes
31..40	medium	no	excellent	yes
31..40	high	yes	fair	yes
>40	medium	no	excellent	no

Output: A Decision Tree for "buys_computer"



Rule Extraction from a Decision Tree

- Rules are easier to understand than large trees
- One rule is created for each path from the root to a leaf
- Each attribute-value pair along a path forms a conjunction: the leaf holds the class prediction
- Rules are mutually exclusive and exhaustive

- Example: Rule extraction from our *buys_computer* decision-tree

IF age <=30 AND student = no THEN buys_computer = no
 IF age <= 30 AND student = yes THEN buys_computer = yes
 IF age = 31...40 THEN buys_computer = yes
 IF age > 40 AND credit_rating = excellent THEN buys_computer = no
 IF age > 40 AND credit_rating = fair THEN buys_computer = yes

23

Gain Ratio for Attribute Selection (C4.5)

- Information gain measure is biased towards attributes with a large number of values
- C4.5 (a successor of ID3) uses gain ratio to overcome the problem (normalization to information gain)

$$\text{SplitInfo}_A(D) = -\sum_{j=1}^n \frac{|D_j|}{|D|} \log_2 \left(\frac{|D_j|}{|D|} \right)$$

- GainRatio(A) = Gain(A)/SplitInfo(A)
- Ex. $\text{SplitInfo}_A(D) = -\frac{4}{14} \times \log_2 \left(\frac{4}{14} \right) - \frac{6}{14} \times \log_2 \left(\frac{6}{14} \right) - \frac{4}{14} \times \log_2 \left(\frac{4}{14} \right) = 0.926$
 - gain_ratio(income) = 0.029/0.926 = 0.031
- The attribute with the maximum gain ratio is selected as the splitting attribute

24

Gini index (CART, IBM IntelligentMiner)

- A measure of impurity to decide how often a randomly chosen element from the set would be incorrectly labeled if it were randomly labeled according to the distribution of labels in the subset.
- If D a dataset containing examples from m classes, then

$$gini(D) = \sum_{i=1}^m p_i(1-p_i) = \sum_{i=1}^m p_i - \sum_{i=1}^m p_i^2 = 1 - \sum_{i=1}^m p_i^2$$
 Where p_i is the relative frequency of class i in D
- If a data set D is split on A into two subsets D_1 and D_2 , the *gini* index $gini_A(D)$ is defined as

$$gini_A(D) = \frac{|D_1|}{|D|} gini(D_1) + \frac{|D_2|}{|D|} gini(D_2)$$
- Reduction in Impurity: $\Delta gini(A) = gini(D) - gini_A(D)$
- The attribute provides the smallest $gini_{split}(D)$ (or the largest reduction in impurity) is chosen to split the node (*need to enumerate all the possible splitting points for each attribute*)

25

Example: Gini index

- Ex. D has 9 tuples in `buys_computer` = "yes" and 5 in "no"

$$gini(D) = 1 - \left(\frac{9}{14}\right)^2 - \left(\frac{5}{14}\right)^2 = 0.459$$
- Suppose the attribute `income` partitions D into 10 in D_1 : {low, medium} and 4 in D_2

$$gini_{income\{low,medium\}}(D) = \left(\frac{10}{14}\right) gini(D_1) + \left(\frac{4}{14}\right) gini(D_2)$$

$$= \frac{10}{14} \left(1 - \left(\frac{6}{10}\right)^2 - \left(\frac{4}{10}\right)^2\right) + \frac{4}{14} \left(1 - \left(\frac{1}{4}\right)^2 - \left(\frac{3}{4}\right)^2\right)$$

$$= 0.450$$

$$= gini_{income \in \{high\}}(D)$$
 but $gini_{(medium,high)}$ is 0.30 and thus the best since it is the lowest
- All attributes are assumed continuous-valued
- May need other tools, e.g., clustering, to get the possible split values
- Can be modified for categorical attributes

26

Comparing Attribute Selection Measures

- The three measures, in general, return good results but
 - Information gain:
 - biased towards multivalued attributes
 - Gain ratio:
 - tends to prefer unbalanced splits in which one partition is much smaller than the others
 - Gini index:
 - biased to multivalued attributes
 - has difficulty when # of classes is large
 - tends to favor tests that result in equal-sized partitions and purity in both partitions

27

Overfitting and Tree Pruning

- Overfitting:** Occurs when a model describes random error or noise instead of the underlying relationship.
- An induced tree may overfit the training data
 - Too many branches, some may reflect anomalies due to noise or outliers
 - Poor accuracy for unseen samples
- Two approaches to avoid overfitting
 - Prepruning:** Halt tree construction early—do not split a node if this would result in the goodness measure falling below a threshold
 - Difficult to choose an appropriate threshold
 - Postpruning:** Remove branches from a "fully grown" tree - get a sequence of progressively pruned trees
 - Use a set of data different from the training data to decide which is the "best pruned tree"

28

Pros & Cons of Using DT

- Pros:**
 - Simple to understand and interpret
 - Requires little data preparation
 - Able to handle both numerical and categorical data
 - Robust
 - Perform well with large data in short time
- Cons:**
 - Learning an optimal decision tree is NP-complete
 - Decision-tree learners create over-complex trees that do not generalize the data well.

29

Bayesian Classification

- A statistical classifier:** predicts class membership probabilities
- Foundation:** Based on Bayes' Theorem.
- Performance:** Comparable performance with decision tree and selected neural network classifiers
- Incremental:** Each training example can incrementally increase/decrease the probability that a hypothesis is correct
 - prior knowledge can be combined with observed data

Bayesian Theorem: Basics

- Let \mathbf{X} be a data sample ("evidence"): class label is unknown
- Let H be a hypothesis that \mathbf{X} belongs to class C
- Classification is to determine $P(H|\mathbf{X})$, the probability that the hypothesis holds given the observed data sample \mathbf{X}
- $P(H)$ - prior probability
 - e.g., \mathbf{X} will buy computer, regardless of age, income, ...
- $P(\mathbf{X})$: probability that sample data is observed
- $P(\mathbf{X}|H)$ (posteriori probability), the probability of observing the sample \mathbf{X} , given that the hypothesis H holds
 - e.g., Given that \mathbf{X} will buy computer, the prob. that \mathbf{X} is 31..40, medium income

Bayesian Theorem

- Given training data \mathbf{X} , posteriori probability of a hypothesis H , $P(H|\mathbf{X})$, follows the Bayes theorem

$$P(H|\mathbf{X}) = \frac{P(\mathbf{X}|H)P(H)}{P(\mathbf{X})}$$

- Informally, this can be written as
posteriori = likelihood \times prior / evidence
- Predicts \mathbf{X} belongs to C_i iff the probability $P(C_i|\mathbf{X})$ is the highest among all the $P(C_k|\mathbf{X})$ for all the k classes
- Practical difficulty:** require initial knowledge of many probabilities, significant computational cost

Naïve Bayesian Classifier

- Let \mathbf{D} be a training set of tuples and their associated class labels, and each tuple is represented by an n -D attribute vector $\mathbf{X} = (x_1, x_2, \dots, x_n)$
- Suppose there are m classes C_1, C_2, \dots, C_m
- Classification is to derive the maximum posteriori, i.e., the maximal $P(C_i|\mathbf{X})$
- This can be derived from Bayes' theorem $P(C_i|\mathbf{X}) = \frac{P(\mathbf{X}|C_i)P(C_i)}{P(\mathbf{X})}$
- Since $P(\mathbf{X})$ is constant for all classes, only $P(C_i|\mathbf{X}) = P(\mathbf{X}|C_i)P(C_i)$ needs to be maximized

Derivation of Naïve Bayes Classifier

- Assumption:** attributes are conditionally independent

$$P(\mathbf{X}|C_i) = \prod_{k=1}^n P(x_k | C_i) = P(x_1 | C_i) \times P(x_2 | C_i) \times \dots \times P(x_n | C_i)$$

- This greatly reduces the computation cost: Only counts the class distribution
- If A_k is categorical, $P(x_k | C_i)$ is the # of tuples in C_i having value x_k for A_k divided by $|C_i|$ (# of tuples of C_i in \mathbf{D})
- If A_k is continuous-valued, $P(x_k | C_i)$ is usually computed based on Gaussian distribution with a mean μ and standard deviation σ

$$g(x, \mu, \sigma) = \frac{1}{\sqrt{2\pi}\sigma} e^{-\frac{(x-\mu)^2}{2\sigma^2}}$$

and $P(x_k | C_i)$ is $P(\mathbf{X}|C_i) = g(x_k, \mu_{C_i}, \sigma_{C_i})$

Naïve Bayesian Classifier: Training Dataset

Class:

C1: buys_computer = 'yes'

C2: buys_computer = 'no'

Data sample:

$\mathbf{X} = (\text{age} \leq 30, \text{Income} = \text{medium}, \text{Student} = \text{yes}, \text{Credit_rating} = \text{Fair})$

	age	income	student	credit_rating	comp
	<=30	high	no	fair	no
	<=30	high	no	excellent	no
	31...40	high	no	fair	yes
	>40	medium	no	fair	yes
	>40	low	yes	fair	yes
	>40	low	yes	excellent	no
	31...40	low	yes	excellent	yes
	<=30	medium	no	fair	no
	<=30	low	yes	fair	yes
	>40	medium	yes	fair	yes
	<=30	medium	yes	excellent	yes
	31...40	medium	no	excellent	yes
	31...40	high	yes	fair	yes
	>40	medium	no	excellent	no

Naïve Bayesian Classifier: An Example

- $P(\text{buys_computer} = \text{"yes"}) = 9/14 = 0.643$
 $P(\text{buys_computer} = \text{"no"}) = 5/14 = 0.357$

- Compute $P(\mathbf{X}|C_i)$ for each class

$P(\text{age} = \text{"<=30"} | \text{buys_computer} = \text{"yes"}) = 2/9 = 0.222$
 $P(\text{age} = \text{"<=30"} | \text{buys_computer} = \text{"no"}) = 3/5 = 0.6$
 $P(\text{income} = \text{"medium"} | \text{buys_computer} = \text{"yes"}) = 4/9 = 0.444$
 $P(\text{income} = \text{"medium"} | \text{buys_computer} = \text{"no"}) = 2/5 = 0.4$
 $P(\text{student} = \text{"yes"} | \text{buys_computer} = \text{"yes"}) = 6/9 = 0.667$
 $P(\text{student} = \text{"yes"} | \text{buys_computer} = \text{"no"}) = 1/5 = 0.2$
 $P(\text{credit_rating} = \text{"fair"} | \text{buys_computer} = \text{"yes"}) = 6/9 = 0.667$
 $P(\text{credit_rating} = \text{"fair"} | \text{buys_computer} = \text{"no"}) = 2/5 = 0.4$

$\mathbf{X} = (\text{age} \leq 30, \text{income} = \text{medium}, \text{student} = \text{yes}, \text{credit_rating} = \text{fair})$

$P(\mathbf{X}|\text{buys_computer} = \text{"yes"}) = 0.222 \times 0.444 \times 0.667 \times 0.667 = 0.044$
 $P(\mathbf{X}|\text{buys_computer} = \text{"no"}) = 0.6 \times 0.4 \times 0.2 \times 0.4 = 0.019$
 $P(\text{buys_computer} = \text{"yes"}|\mathbf{X}) = P(\mathbf{X}|\text{buys_computer} = \text{"yes"}) \times P(\text{buys_computer} = \text{"yes"}) = 0.028$
 $P(\mathbf{X}|\text{buys_computer} = \text{"no"}|\mathbf{X}) = P(\mathbf{X}|\text{buys_computer} = \text{"no"}) \times P(\text{buys_computer} = \text{"no"}) = 0.007$

Therefore, \mathbf{X} belongs to class ("buys_computer = yes")

Avoiding the 0-Probability Problem


- Naïve Bayesian prediction requires each conditional prob. be non-zero. Otherwise, the predicted prob. will be zero
- Suppose a dataset with 1000 tuples, income=low (0), income= medium (990), and income = high (10)
- Use **Laplacian correction** (or Laplacian estimator)
 - Adding 1 to each case
 - $\text{Prob}(\text{income} = \text{low}) = 1/1003$
 - $\text{Prob}(\text{income} = \text{medium}) = 991/1003$
 - $\text{Prob}(\text{income} = \text{high}) = 11/1003$

Naïve Bayesian Classifier: Comments

- Advantages**
 - Easy to implement
 - Good results obtained in most of the cases
- Disadvantages**
 - Assumption: class conditional independence, therefore loss of accuracy
 - Practically, dependencies exist among variables
 - e.g., **Symptoms**: fever, cough etc., **Disease**: lung cancer, diabetes, etc.
 - Dependencies among these cannot be modeled by Naïve Bayesian Classifier
- How to deal with these dependencies?**
 - Bayesian Belief Networks

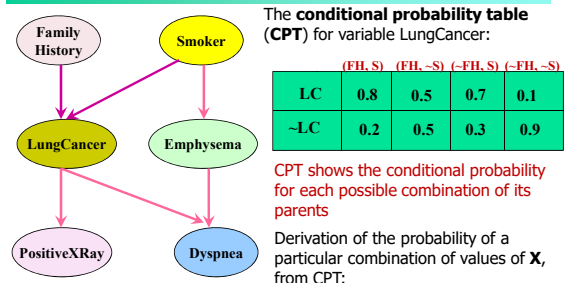
Bayesian Belief Networks

- Bayesian belief network allows a *subset* of the variables conditionally independent
- A graphical model of causal relationships
 - Represents dependency among the variables
 - Gives a specification of joint probability distribution



- Nodes**: random variables
- Links**: dependency
- X and Y are the parents of Z, and Y is the parent of P
- No dependency between Z and P
- Has no loops or cycles

Bayesian Belief Network: An Example



Bayesian Belief Networks

$$P(x_1, \dots, x_n) = \prod_{i=1}^n P(x_i \mid \text{Parents}(Y_i))$$

Lazy vs. Eager Learning

- Lazy learning**: Simply stores training data (or only minor processing) and waits until it is given a test tuple.
- Eager learning**: Given a set of training set, constructs a classification model before receiving new (e.g., test) data to classify
- Pros & Cons**:
 - Lazy requires less time in training but more time in predicting
 - Lazy methods are more accurate as they use many local linear functions to form its implicit global approximation to the target function. Whereas, eager methods must commit to a single hypothesis that covers the entire instance space.

Evaluation: What is Good Classification?

- Correct classification**: The known label of test sample is identical with the **class result** from the classification model
- Accuracy ratio**: the percentage of test set samples that are correctly classified by the model
- A **distance measure** between classes can be used
 - e.g., classifying "football" document as a "basketball" document is not as bad as classifying it as "crime".

Accuracy?

Accuracy is the percentage of test set samples that are correctly classified by the model

Given m classes, CM_{ij} , an entry in a **confusion matrix**, indicates # of tuples in class i that are labeled by the classifier as class j

Confusion Matrix

		System Classification	
		Class-1 (+ve)	Class-2 (-ve)
Actual Classification	Class-1 (+ve)	True Positive (TP)	False Negative (FN)
	Class-2 (-ve)	False Positive (FP)	True Negative (TN)

$$accuracy = \frac{TP + TN}{TP + TN + FP + FN}$$

Error rate (misclassification rate) = $1 - accuracy$

Other Measures

- Alternative accuracy measures (e.g., for cancer diagnosis)

$$sensitivity = \frac{TP}{P} = \frac{TP}{TP + FN} \quad /* \text{ true positive recognition rate } */$$

$$specificity = \frac{TN}{N} = \frac{TN}{TN + FP} \quad /* \text{ true negative recognition rate } */$$

$$precision = \frac{TP}{TP + FP}$$

$$recall = \frac{TP}{TP + FN}$$

$$accuracy = sensitivity \times \frac{P}{P + N} + specificity \times \frac{N}{P + N}$$

/* P = total positives = $TP + FN$; N = total negatives = $TN + FP$ */

Predictor Error Measures

- Measure predictor accuracy:** measure how far off the predicted value is from the actual known value
- Loss function:** measures the error between y_i and the predicted value y_i'
 - Absolute error:** $|y_i - y_i'|$
 - Squared error:** $(y_i - y_i')^2$
- Test error (generalization error):** the average loss over the test set

$$mean \ absolute \ error = \frac{\sum_{i=1}^d |y_i - y_i'|}{d} \quad mean \ squared \ error = \frac{\sum_{i=1}^d (y_i - y_i')^2}{d}$$

$$relative \ absolute \ error = \frac{\sum_{i=1}^d |y_i - y_i'|}{\sum_{i=1}^d |y_i - \bar{y}|} \quad relative \ squared \ error = \frac{\sum_{i=1}^d (y_i - y_i')^2}{\sum_{i=1}^d (y_i - \bar{y})^2}$$

The mean squared-error exaggerates the presence of outliers

Evaluating the Accuracy of a Classifier or Predictor

- Holdout method**
 - Given data is randomly partitioned into two independent sets
 - Training set (e.g., 2/3) for model construction
 - Test set (e.g., 1/3) for accuracy estimation
- Random sampling** (a variation of holdout)
 - Repeat holdout k times, accuracy = avg. of the accuracies obtained
- Cross-validation** (k -fold, where $k = 10$ is most popular)
 - Randomly partition the data into k mutually exclusive subsets, each approximately equal size
 - At i -th iteration, use D_i as test set and others as training set
 - Leave-one-out:** k folds where $k = \#$ of tuples, for small sized data
 - Stratified cross-validation:** folds are stratified so that class dist. in each fold is approx. the same as that in the initial data

Cont...

- Bootstrap**
 - Samples the given training tuples uniformly with replacement
 - i.e., Each time a tuple is selected, it is equally likely to be selected again and readded to the training set.
 - 0.632 bootstrap:** A commonly used bootstrap method
 - Given a dataset of d tuples, it is sampled d times, with replacement, resulting in a bootstrap sample or training set of d samples.

$$prob(selection) = \frac{1}{d}$$

$$prob(no \ selection) = 1 - \frac{1}{d}$$

$$In \ d \ repetitions, \ prob(no \ selection) = \left(1 - \frac{1}{d}\right)^d \approx e^{-1} = 0.368$$

If d is very large

36.8% of tuples will not be selected for training set – used for test

Remaining 63.2% tuples will form the training set

Cont...


- The sampling procedure in bootstrap can be repeated k times
- The overall accuracy of model M is then calculated as

$$Acc(M) = \sum_{i=1}^k (0.632 \times Acc(M_i)_{test_set} + 0.368 \times Acc(M_i)_{train_set})$$

- Works well for small data sets

Ensemble Methods: Increasing the Accuracy

- **Ensemble methods**
 - Methods that use a combination of models to increase accuracy
 - Combine a series of k learned models, M_1, M_2, \dots, M_k , with the aim of creating an improved model M^*
 - Can work for both **classification and prediction**



49

Popular Ensemble Methods

- **Bagging:**
 - Using majority vote for classification
 - Averaging the predicted values for prediction
- **Boosting:**
 - Assigning weights to classifiers and using weighted vote for classification and prediction

April 6, 2019

Data Mining: Concepts and Techniques

50

Bagging

- **Analogy:** Diagnosis based on multiple doctors' majority vote
- **Training:** learning models
 - Given a set D of d tuples, at each iteration i , a training set D_i of d tuples is sampled with replacement from D (i.e., bootstrap)
 - A classifier model M_i is learned for each training set D_i
- **Classification:** classifying an unknown sample X
 - Each classifier M_i returns its class prediction
 - The bagged classifier M^* counts the votes and assigns the class with the most votes to X
- **Prediction:** can be applied to the prediction of continuous values by taking the average value of each prediction for a given test tuple

51

Boosting

- **Analogy:** Consult several doctors, based on a combination of weighted diagnoses – weight assigned based on the previous diagnosis accuracy
- **How boosting works?**
 - Weights are assigned to each training tuple
 - A series of k classifiers is iteratively learned
 - After a classifier M_i is learned, the weights are updated to allow the subsequent classifier, M_{i+1} , to pay more attention to the training tuples that were misclassified by M_i
 - The final M^* combines the votes of each individual classifier, where the weight of each classifier's vote is a function of its accuracy
- Boosting tends to **achieve greater accuracy than bagging**, but it also risks Overfitting the model to misclassified data

52

Adaboost (Freund and Schapire, 1997)

- A popular boosting algorithm
- Given a set of d class-labeled tuples, $(X_1, y_1), \dots, (X_d, y_d)$
- Initially, all the weights of tuples are set the same ($1/d$)
- Generate k classifiers in k rounds. At round i ,
 - Tuples from D are sampled (with replacement) to form a training set D_i of the same size
 - Each tuple's chance of being selected is based on its weight
 - A classification model M_i is derived from D_i
 - Its error rate is calculated using D_i as a test set
 - If a tuple is correctly classified, its weight is decreased
 - If a tuple is incorrectly classified, its weight is increased

53

Adaboost (Cont...)

- A tuple's weight reflects how hard it is to classify
 - The higher the weight, the more often it has been misclassified.

54

Computing Error Rate of a Model

- Error rate of classifier M_i is the sum of the weights of the misclassified tuples in D_i , that is

$$error(M_i) = \sum_{j=1}^d w_j \times error(X_j)$$

Where, $error(X_j)$ is the misclassification error of tuple X_j (if misclassification then 1, otherwise 0).

- If the performance of classifier M_i is so poor that its error exceeds 0.5, then we abandon it, and repeat the sampling to re-learn model M_i

55

Tuples Weight Updation

- If a tuple in round i was correctly classified, its weight is multiplied by

$$\frac{error(M_i)}{1 - error(M_i)}$$

- The weights for all tuples (including the misclassified ones) are normalized so that their sum remains the same as it was before.
- To normalize a weight, multiply it by the sum of the old weights, and divided by the sum of the new weights.
- As a result, the weights of misclassified tuples are increased and the weights of correctly classified tuples are decreased

56

X	X/(1-X)
0.10	0.11
0.20	0.25
0.30	0.43
0.40	0.67
0.50	1.00
0.60	1.50
0.70	2.33
0.80	4.00
0.90	9.00

57

Initial weight	prediction	Error Calculation	Weight updation	Weight Normalization
0.10	T	0.00	0.07	0.03
0.10	F	0.10	0.10	0.05
0.10	F	0.10	0.10	0.05
0.10	T	0.00	0.07	0.03
0.10	T	0.00	0.07	0.03
0.10	F	0.10	0.10	0.05
0.10	T	0.00	0.07	0.03
0.10	F	0.10	0.10	0.05
0.10	T	0.00	0.07	0.03
0.10	T	0.00	0.07	0.03
Error(Mi)		0.40	0.80	0.40
Multiplication factor		0.67		

58

Using Classifier Weights for Prediction

- The lower a classifier's error rate, the more accurate it is, and therefore, the higher its weight for voting should be. The weight of classifier M_i 's vote is $\log \frac{1 - error(M_i)}{error(M_i)}$
- For each class, c , we sum the weights of each classifier that assigned class c to X . The class with the highest sum is the "winner" and is returned as the class prediction for tuple X .

59

How Adaboost learns models?

Steps:

- (1) initialize the weight of each tuple in D to $1/d$;
- (2) for $i = 1$ to k do // for each round:
 - (3) sample D with replacement according to the tuple weights to obtain D_i ;
 - (4) use training set D_i to derive a model, M_i ;
 - (5) compute $error(M_i)$, the error rate of M_i (Equation 6.66)
 - (6) if $error(M_i) > 0.5$ then
 - (7) reinitialize the weights to $1/d$
 - (8) go back to step 3 and try again;
 - (9) endif
 - (10) for each tuple in D_i that was correctly classified do
 - (11) multiply the weight of the tuple by $error(M_i)/(1 - error(M_i))$; // update weights
 - (12) normalize the weight of each tuple;
 - (13) endfor

60

How Adaboost classifies a tuple X?

- **Steps:**
 - Initialize weight of each class to 0
 - For $i=1$ to k do // for each classifier
 - $w_i = \log((1 - \text{error}(M_i)) / \text{error}(M_i))$ // weight of the classifier's vote
 - $c = M_i(X)$ // get class prediction for X from M_i
 - Add w_i to weight for class c
 - Endfor
 - Return the class with the largest weight

61


Model Selection: ROC Curves

- Receiver Operating Characteristics (ROC) curves are used for **visual comparison** of classification models
- Originated from **signal detection theory** developed during World War II for analysis of Radar images
- Presents the trade-off between the **true positive rate** and the **false positive rate**
- The area under the ROC curve is a **measure of the accuracy** of the model
- The closer to the diagonal line (i.e., the closer the area is to 0.5), the less accurate is the model
- A model with **perfect accuracy** will have an area of 1.0

62

How to Draw ROC Curve?


- **Rank the test tuples in decreasing order:** the one that is most likely to belong to the positive class appears at the top of the list
- Vertical axis represents the true positive rate
- Horizontal axis represents the false positive rate
- The plot also shows a diagonal line



63

ROC Curve Plotting as a Step Function

Inst#	Class	Score	Inst#	Class	Score
1	p	.9	11	p	.4
2	p	.8	12	n	.39
3	n	.7	13	p	.38
4	p	.6	14	n	.37
5	p	.55	15	n	.36
6	p	.54	16	n	.35
7	n	.53	17	p	.34
8	n	.52	18	n	.33
9	p	.51	19	p	.30
10	n	.505	20	n	.1



References

- C. Apte and S. Weiss. **Data mining with decision trees and decision rules**. Future Generation Computer Systems, 13, 1997.
- C. M. Bishop, **Neural Networks for Pattern Recognition**. Oxford University Press, 1995.
- L. Breiman, J. Friedman, R. Olshen, and C. Stone. **Classification and Regression Trees**. Wadsworth International Group, 1984.
- C. J. C. Burges. **A Tutorial on Support Vector Machines for Pattern Recognition**. *Data Mining and Knowledge Discovery*, 2(2): 121-168, 1998.
- P. K. Chan and S. J. Stolfo. **Learning arbiter and combiner trees from partitioned data for scaling machine learning**. KDD'95.
- W. Cohen. **Fast effective rule induction**. ICML'95.
- G. Cong, K.-L. Tan, A. K. H. Tung, and X. Xu. **Mining top-k covering rule groups for gene expression data**. SIGMOD'05.
- A. J. Dobson. **An Introduction to Generalized Linear Models**. Chapman and Hall, 1990.
- G. Dong and J. Li. **Efficient mining of emerging patterns: Discovering trends and differences**. KDD'99.

April 6, 2019

Data Mining: Concepts and Techniques

65

References

- R. O. Duda, P. E. Hart, and D. G. Stork. **Pattern Classification**, 2ed. John Wiley and Sons, 2001.
- U. M. Fayyad. **Branching on attribute values in decision tree generation**. AAAI' 94.
- Y. Freund and R. E. Schapire. **A decision-theoretic generalization of on-line learning and an application to boosting**. J. Computer and System Sciences, 1997.
- J. Gehrke, R. Ramakrishnan, and V. Ganti. **Rainforest: A framework for fast decision tree construction of large datasets**. VLDB' 98.
- J. Gehrke, V. Ganti, R. Ramakrishnan, and W.-Y. Loh. **BOAT -- Optimistic Decision Tree Construction**. SIGMOD'99.
- T. Hastie, R. Tibshirani, and J. Friedman. **The Elements of Statistical Learning: Data Mining, Inference, and Prediction**. Springer-Verlag, 2001.
- D. Heckerman, D. Geiger, and D. M. Chickering. **Learning Bayesian networks: The combination of knowledge and statistical data**. Machine Learning, 1995.
- M. Kamber, L. Winstone, W. Gong, S. Cheng, and J. Han. **Generalization and decision tree induction: Efficient classification in data mining**. RIDE'97.
- B. Liu, W. Hsu, and Y. Ma. **Integrating Classification and Association Rule**. KDD'98.
- W. Li, J. Han, and J. Pei. **CMAR: Accurate and Efficient Classification Based on Multiple Class-Association Rules**. ICDM'01.

April 6, 2019

Data Mining: Concepts and Techniques

66

References

- T.-S. Lim, W.-Y. Loh, and Y.-S. Shih. **A comparison of prediction accuracy, complexity, and training time of thirty-three old and new classification algorithms.** *Machine Learning*, 2000.
- J. Magidson. **The Chaid approach to segmentation modeling: Chi-squared automatic interaction detection.** In R. P. Bagozzi, editor, *Advanced Methods of Marketing Research*, Blackwell Business, 1994.
- M. Mehta, R. Agrawal, and J. Rissanen. **SLIQ : A fast scalable classifier for data mining.** *EDBT'96*.
- T. M. Mitchell. **Machine Learning.** McGraw Hill, 1997.
- S. K. Murthy, **Automatic Construction of Decision Trees from Data: A Multi-Disciplinary Survey**, *Data Mining and Knowledge Discovery* 2(4): 345-389, 1998
- J. R. Quinlan. **Induction of decision trees.** *Machine Learning*, 1:81-106, 1986.
- J. R. Quinlan and R. M. Cameron-Jones. **FOIL: A midterm report.** *ECML'93*.
- J. R. Quinlan. **C4.5: Programs for Machine Learning.** Morgan Kaufmann, 1993.
- J. R. Quinlan. **Bagging, boosting, and c4.5.** *AAAI'96*.

April 6, 2019

Data Mining: Concepts and Techniques

67

References

- R. Rastogi and K. Shim. **Public: A decision tree classifier that integrates building and pruning.** *VLDB'98*.
- J. Shafer, R. Agrawal, and M. Mehta. **SPRINT : A scalable parallel classifier for data mining.** *VLDB'96*.
- J. W. Shavlik and T. G. Dietterich. **Readings in Machine Learning.** Morgan Kaufmann, 1990.
- P. Tan, M. Steinbach, and V. Kumar. **Introduction to Data Mining.** Addison Wesley, 2005.
- S. M. Weiss and C. A. Kulikowski. **Computer Systems that Learn: Classification and Prediction Methods from Statistics, Neural Nets, Machine Learning, and Expert Systems.** Morgan Kaufman, 1991.
- S. M. Weiss and N. Indurkha. **Predictive Data Mining.** Morgan Kaufmann, 1997.
- I. H. Witten and E. Frank. **Data Mining: Practical Machine Learning Tools and Techniques**, 2ed. Morgan Kaufmann, 2005.
- X. Yin and J. Han. **CPAR: Classification based on predictive association rules.** *SDM'03*
- H. Yu, J. Yang, and J. Han. **Classifying large data sets using SVM with hierarchical clusters.** *KDD'03*.

April 6, 2019

Data Mining: Concepts and Techniques

68